

Performance Analysis and Optimization

Raunak Borker

Part1 Implementation Time Comparison

- Performance comparison over **25 steps**
- **Python** implementation: **7.1sec**
- **C++** implementation: **0.3sec**
- **24x** speedup! (not good enough)

Expectation before analyzing

- The **bottleneck** in performance could probably be due to **InternalForce** computations as it involves N computations for every element - total of **N^2** computations

Profiling and Analysis

- gprofile results indicated **ComputeEnergy**, **InternalForce** and **PutInBox** methods are major time consumers
- **std::vector append** also showed up as one of the heavy users
- cachegrind indicated **~0.0% L1 miss rates**

Part 1 Implementation

- Particle class (single particle)
 - Vectorial quantities like position, velocity, acceleration defined as **std::vector**
- ParticleBox class (collection of particles in a box)
 - Methods (for e.g. InternalForce, ComputeEnergy, Displacement) all returned **std::vector** as output

Part 2 Implementation

- Particle class (single particle)
 - Vectorial quantities now defined as **std::array** replacing **std::vector**, since quantities are of a fixed size (3 X 1)
- ParticleBox class (collection of particles in a box)
 - Methods mentioned previously now return void, but accept an additional input argument as a reference to **std::array** which stores the output

Part 2 Implementation

- L1 miss rate was good in Part 1
- Part 2 - 2.5% miss rate

```
==3665== I    refs:      1,458,807
==3665== I1   misses:      1,394
==3665== LLi  misses:      1,364
==3665== I1   miss rate:    0.09%
==3665== LLi  miss rate:    0.09%
==3665==
==3665== D    refs:      485,655 (357,939 rd + 127,716 wr)
==3665== D1   misses:      12,327 ( 10,400 rd +   1,927 wr)
==3665== LLd  misses:       7,676 (   6,170 rd +   1,506 wr)
==3665== D1   miss rate:    2.5% (   2.9% +   1.5% )
==3665== LLd  miss rate:    1.5% (   1.7% +   1.1% )
```

Performance Comparison

- Performance comparison over **25 steps**
 - **Python** implementation: **7.1sec**
 - **C++** implementation: **0.038sec**
 - **187x** speedup!

gprofile output (100 steps)

% time	cumulative seconds	self seconds	calls	self us/call	total us/call	name
100.08	0.05	0.05	12800	3.91	3.91	ParticleBox<double>::InternalForce(unsigned int)
0.00	0.05	0.00	1400	0.00	0.00	void std::vector<std::array<double, 3ul>, std::allocator<std::array<double, 3ul> > >:: M_emplace_back_aux<std::array<double, 3ul> >(std::array<double, 3ul>&&)
0.00	0.05	0.00	101	0.00	0.00	ParticleBox<double>::ComputeEnergy(std::array<double, 3ul>&)
0.00	0.05	0.00	1	0.00	0.00	ParticleBox<double>::ParticleBox(std::string)
0.00	0.05	0.00	1	0.00	0.00	main

- As can be seen the initial expectation of **InternalForce** being bottleneck is now revealed
- This could possibly be due since this operation is not “vectorized”, i.e. the force contribution to particle i due to j is equal and opposite, but in the current implementation is computed twice!